

# Implementing Chib (1998) multiple change-point model

## Anh Le

April 30, 2015

The `line_profiler` extension is already loaded. To reload it, use:  
`%reload_ext line_profiler`  
The `autoreload` extension is already loaded. To reload it, use:  
`%reload_ext autoreload`

Github repo: [https://github.com/LaDilettante/changepoint\\_model](https://github.com/LaDilettante/changepoint_model)

## 1 Introduction

This project implements Chib (1998) multiple change-point model. This algorithm comprises of two parts:  
Given a number of change-points 1. Detect the location of change-points and the parameters of the DGP in the different regimes 2. Calculate the marginal likelihood of the model, which can be used to compare models with different change-points

The innovation of the model is “a formulation of the change-point model in terms of a latent discrete statement variable that indicates the regime from which a particular observation has been drawn.” Formulating the model in this way, we can then use the forward-backward algorithm to sample the latent state vector (and Gibbs sampler to sample the other parameters).

This method requires only  $2n$  passes through the data, *regardless of the number of change-points*. This is a big improvement over previous change-point models, in which we have to sample the change-point one at a time.

## 2 Model Parameterization

Consider a time series  $Y_n$  with an unknown  $m$  change-points.

Introduce the latent state vector  $s_t = \{1, 2, \dots, m+1\}$ , indicating which regime the observation  $y_t$  is drawn from. In other words, if  $s_t = k$ , then  $y_t \sim f(y_t | Y_{t-1}, \theta_k)$

The variable  $s_t$  is modeled as a discrete Markov process with the transition probability matrix designed to fit the change-point model. Specifically,

$$P = \begin{pmatrix} p_{11} & p_{22} & 0 & \dots & 0 \\ 0 & p_{22} & p_{23} & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ \dots & \vdots & 0 & p_{mm} & p_{m,m+1} \\ 0 & 0 & 0 & \dots & 1 \end{pmatrix}$$

Given this transition matrix, the latent state variable  $s_t$  can only jump forward, conforming to the change-point model.

**In sum**, we have 3 following parameters  $S_n$ ,  $\Theta$ ,  $P$ . The next section describe the MCMC scheme to sample each parameter.

### 3 MCMC scheme

With a prior density  $\pi(\Theta, P)$  and data  $Y_n$ , we want to estimate the posterior  $\pi(\Theta, P|Y_n)$ . With the data augmentation using the latent variable  $s_t$ , we can use the following full conditionals:

- $S_n|Y_n, \Theta, P$
- $P|S_n$
- $\Theta|Y_n, S_n$

Note how the introduction of  $S_n$  simplifies the full conditional of  $P$  and  $\Theta$

#### 3.1 Simulation of $s_t$ (forward-backward)

Denote  $S_t$  the state history up to time  $t$ , and  $S^{t+1}$  the future from  $t+1$  to  $n$ . We now sample  $p(S_n|Y_n, \Theta, P)$

$$p(S_n|Y_n, \Theta, P) = p(s_{n-1}|Y_n, s_n, \Theta, P) \times \quad (1)$$

$$p(s_{n-2}|Y_n, S^{n-1}, \Theta, P) \times \quad (2)$$

$$\dots \quad (3)$$

$$p(s_1|Y_n, S^2, \Theta, P) \quad (4)$$

A typical term from above is in the form  $p(s_t|Y_n, S^{t+1}, \Theta, P)$ , and can be calculated as following:

$$p(s_t|Y_n, S^{t+1}, \Theta, P) \propto p(s_t|Y_t, \Theta, P)p(s_{t+1}|s_t, P) \quad (5)$$

**Key idea** - Notice how given  $s_{t+1}$  and  $p(s_t|Y_t, \Theta, P)$ , we can calculate  $p(s_t|Y_n, S^{t+1}, \Theta, P)$ . This is how we go **backward** after initializing  $s_n = m+1$

We still need  $p(s_t|Y_t, \Theta, P)$ , which is calculated as:

$$p(s_t = k|Y_t, \Theta, P) = \frac{p(s_t = k|Y_{t-1}, \Theta, P) \times f(y_t|Y_{t-1}, \theta_k)}{\sum_{l=\text{all } k\text{'s}} p(s_t = l|Y_{t-1}, \Theta, P) \times f(y_t|Y_{t-1}, \theta_l)}$$

and

$$p(s_t = k|Y_{t-1}, \Theta, P) = \sum_{l=k-1}^k p_{lk} \times p(s_{t-1} = l|Y_{t-1}, \Theta, P)$$

**Key idea:** - Notice how given  $p(s_{t-1} = l|Y_{t-1}, \Theta, P)$ , we can calculate  $p(s_t = k|Y_t, \Theta, P)$ . This is how we go **forward**, starting by setting  $p(s_1|Y_0, \theta) = 1$ .

#### 3.2 Simulation of $P$

Since  $P|Y_n, S_n, \Theta$  is independent of  $Y_n, \Theta$  given  $S_n$ , sampling from the full conditional is easy.

Given the prior  $p_{ii} \sim \text{Beta}(a, b)$ , we have the posterior  $p_{ii}|S_n \sim \text{Beta}(a + n_{ii}, b + 1)$ , with  $n_{ii}$  denote the number of one-step transitions from state  $i$  to state  $i$  in the sequence  $S_n$

#### 3.3 Simulation of $\Theta$

The full conditional  $\Theta|Y_n, S_n$  depends on each data model. Given the latent state vector  $S_n$ , we have a complete data likelihood, and the posterior update is routine.

## 4 MCEM and calculating marginal likelihood

To decide how many change points there are in the data, we can fit models with different change points the compare the marginal likelihood of each model.

The marginal likelihood of model  $M_r$  is:

$$m(Y_n|M_r) = \frac{f(Y_n|M_r, \Theta^*, P^*)\pi(\Theta^*, P^*|M_r)}{\pi(\Theta^*, P^*|Y_n, M_r)}$$

where  $\Theta^*, P^*$  could be any point in the parameter space. However, for numerical stability, we choose a high posterior density point, which is the MLE, estimated via Monte Carlo EM.

The posterior ordinate  $\pi(\Theta^*, P^*|M_r)$  is separated into two terms:

$$\pi(\Theta^*, P^*|M_r) = \pi(\Theta^*|Y_n)\pi(P^*|Y_n, \Theta^*)$$

both of which are estimated from MCMC samples.

## 5 Implementation

- Each full conditional in a MCMC iteration is a function, grouped into a module (“full\_conditionals.py”). This makes it easy to re-use the sampling step for the MCMC and MCEM
- The update for each parameter in the MCEM iteration is a function, grouped into a module (“mcem.py”)
- The calculation of each term in the marginal likelihood is a function (in file “ordinate.py”)

The MCMC and MCEM samplers then call these modules. To facilitate speed comparison, the sampler can choose to call the regular module or the optimized version (suffixed with “\_opt.py”). The sampler function can also specify the DGP (binary or poisson) and the number of change points.


## 6 Result

We use the algorithm on binary data (simulated) and poisson data (British coal-mining disasters)

### 6.1 Binary data

We simulate binary data  $y_t \sim \text{Bernoulli}(\xi_t)$ , where

$$f(n) = \begin{cases} \theta_1 = 0.5 & \text{if } t \leq 50 \\ \theta_2 = 0.75 & \text{if } 50 < t \leq 100 \\ \theta_3 = 0.25 & \text{if } 100 < t \leq 150 \end{cases}$$

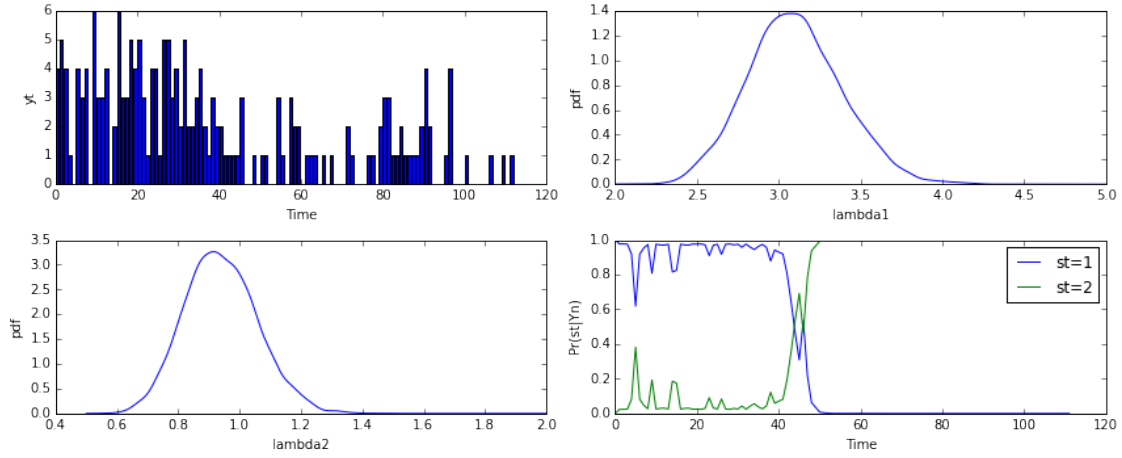


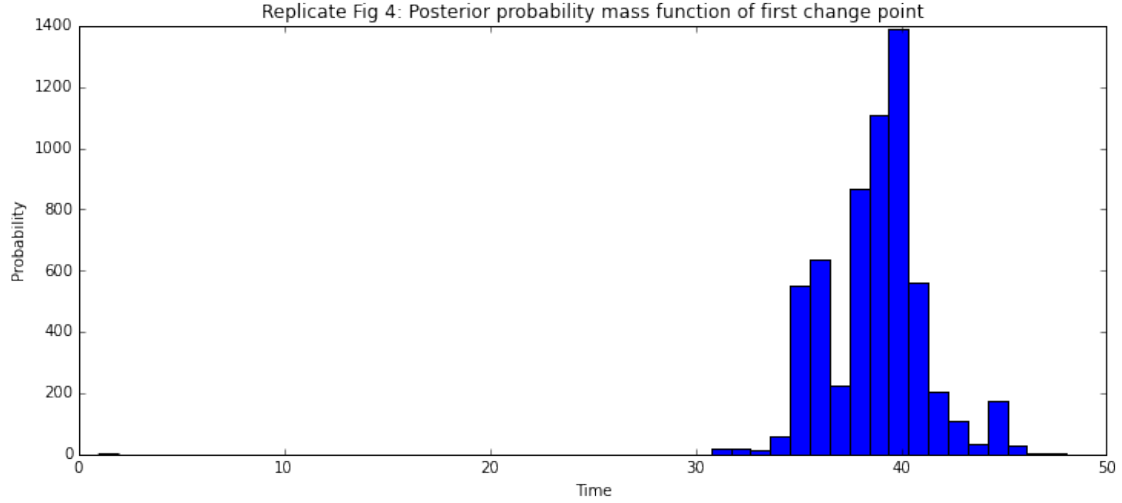
## 6.2 Poisson data

We use the method on the British coal mining disasters by year. Model the count  $y_t$  from a hierarchical Poisson model

$$f(y_t|\xi_t) = \frac{\xi_t^{y_t} e^{-\xi_t}}{y_t!}$$

Our implementation successfully replicates the author's result (Figure 3 and Figure 4) for the model with one break.






### 6.3 Marginal likelihood calculation

The result of our Monte Carlo EM and the marginal likelihood function for Poisson data also matches the author's result exactly. Comparing the marginal likelihood of the one break (-178.37) vs two breaks (-179.59), we reach the same conclusion that one break model is a better fit (See Table 2 in the paper for the same conclusion). The location of change-points in the two-break models also matches the author's exactly.

```
Theta_MLE = [ 3.12513647  0.9260591 ]
log-likelihood = -172.181385853
marginal-likelihood = -178.378503257
```

```
Theta_MLE = [ 3.14525235  1.09034126  0.30937427]
log-likelihood = -170.63154491
marginal-likelihood = -179.592209806
```

Replicate Fig 6: Posterior probability function of the two change points



## 7 Code optimization

Due to the sequential updating in MCMC, we cannot vectorize the entire MCMC step. Instead, we can only vectorize within the sampling of  $\Theta$  and  $P$ . The most time consuming step in each MCMC iteration is the sampling of  $S_n$  – however, due to the nature of the forward-backward algorithm, we cannot speed this part up either.

The following demonstration shows that the vectorize version of *Theta\_sampling* cuts the time in half (from 0.005s to 0.002s). It is important to point out that sampling Theta takes up only a small part of the total time in the MCMC, however.

- Unoptimized

Timer unit: 1e-06 s

Total time: 0.005355 s

File: changepoint/full\_conditionals.py

Function: Theta\_sampling at line 33

Line #	Hits	Time	Per Hit	% Time	Line Contents
33					def Theta_sampling(Yn, Sn, model):
34					'''
35					Sample Theta from its full conditional distribution
36					
37					Args:
38					Yn: n x 1 data vector
39					Sn: n x 1 latent state vector

40					
41					Returns:
42					a (m + 1) x 1 vector of parameter, one param f
43					'''
44	1	48	48.0	0.9	number_of_regimes = len(np.unique(Sn))
45	1	5	5.0	0.1	thetas = np.empty(number_of_regimes)
46					
47	3	8	2.7	0.1	for k in range(1, number_of_regimes + 1):
48	2	5293	2646.5	98.8	thetas[k - 1] = Theta_conditional(k, Yn, Sn, mo
49	1	1	1.0	0.0	return thetas

- Optimized

Timer unit: 1e-06 s

Total time: 0.002913 s

File: changepoint/full\_conditionals\_opt.py

Function: Theta\_sampling at line 32

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
32					def Theta_sampling(Yn, Sn, model):
33					'''
34					Sample Theta from its full conditional distribution
35					
36					Args:
37					Yn: n x 1 data vector
38					Sn: n x 1 latent state vector
39					
40					Returns:
41					a (m + 1) x 1 vector of parameter, one param f
42					'''
43	1	52	52.0	1.8	number_of_regimes = len(np.unique(Sn))
44	1	1	1.0	0.0	m = number_of_regimes - 1
45					
46	1	74	74.0	2.5	Nks = Nk(Sn)
47	1	71	71.0	2.4	Uks = Uk(Yn, Sn)
48					
49	1	2	2.0	0.1	if model == "binary":
50					# Prior
51					a = 2 ; b = 2
52					return st.beta(a + Uks, b + Nks - Uks).rvs()
53	1	1	1.0	0.0	elif model == "poisson":
54					# Different priors based on number of breaks
55	1	0	0.0	0.0	if m == 1:
56	1	1	1.0	0.0	a = 2 ; b = 1
57					elif m == 2:
58					a = 3 ; b = 1
59	1	2711	2711.0	93.1	return st.gamma(a + Uks, scale=(1.0 / (b + Nks

Similar speed-up is achieved in the updating of Theta and P in the Monte Carlo EM. Total time (73.6s -> 67.3s), Theta-M-step (0.63s -> 0.03s). Note how sampling  $S_n$  in the E-step takes up the lion-share of computation time (> 98%)

- Unoptimized

Timer unit: 1e-06 s

Total time: 0.627406 s

File: changepoint/mcem.py

Function: theta\_mstep at line 20

Line #	Hits	Time	Per Hit	% Time	Line Contents
20					def theta_mstep(k, Yn, Sns, model):
21					'''
22					Calculate theta_k, using N samples of Sn
23					'''
24	200	292	1.5	0.0	N = Sns.shape[1]
25					
26	200	532	2.7	0.1	Uks = np.zeros(N)
27	200	294	1.5	0.0	Nks = np.zeros(N)
28					
29	30220	19346	0.6	3.1	for N_ in range(N):
30	30020	308828	10.3	49.2	Uks[N_] = Uk(k, Yn, Sns[:, N_])
31	30020	296273	9.9	47.2	Nks[N_] = Nk(k, Sns[:, N_])
32					
33	200	129	0.6	0.0	if model == "binary":
34					return 1.0 * Uks.sum() / Nks.sum()
35	200	116	0.6	0.0	if model == "poisson":
36	200	1596	8.0	0.3	return 1.0 * Uks.sum() / Nks.sum()

Total time: 73.585 s

File: changepoint/mcem\_sampler.py

Function: mcem\_sampler at line 7

Line #	Hits	Time	Per Hit	% Time	Line Contents
7					def mcem_sampler(Yn, model, m, mcem_module, tol=None):
8					'''
9					Calculate the MLE of Theta and P, using Monte Carlo
10					
11					We use EM because the latent state vector is not o
12					We use Monte Carlo method because the Q function i
13					
14					Args
15					Yn: array, time series data
16					model: string, e.g. "binary", "poisson"
17					m: int, number of change points
18					mcem: module, e.g. mcem, mcem_opt
19					Use the regular of optimized version
20					'''
21					# Initialize
22	1	5	5.0	0.0	n = len(Yn) ; m = m
23	1	162	162.0	0.0	P = init.P(m + 1)
24					
25	1	4	4.0	0.0	P[0, 0] = 0.8 # Paper
26	1	2	2.0	0.0	P[0, 1] = 1 - 0.8
27					
28	1	2	2.0	0.0	if (m == 2) and (model == "poisson"):



29					# Without giving a high chance of staying at s
30					P[1, 1] = 0.9
31					P[1, 2] = 1 - 0.9
32					
33					# Initialize Theta
34	1	2	2.0	0.0	if model == "binary":
35					Theta = np.repeat(0.5, m + 1)
36	1	2	2.0	0.0	elif model == "poisson":
37					#Theta = np.repeat(2, m + 1)
38	1	31	31.0	0.0	Theta = 2 + np.random.rand(m + 1)
39					
40					# N is the number of Sn sample
41					# According to paper, start N = 1 and increases ov
42	1	52	52.0	0.0	Ns = np.linspace(1, 300, 10).astype(np.int64)
43	1	6	6.0	0.0	Thetas = np.empty((m + 1, 100))
44	1	4	4.0	0.0	Ps = np.empty((m + 1, m + 1, 100))
45					
46					# Start 100 MCEM steps
47	1	2	2.0	0.0	i = 0
48	101	125	1.2	0.0	while i < 100:
49	100	174	1.7	0.0	N = Ns[i / 10]
50					
51					# E-step
52	100	72758913	727589.1	98.9	Sns = mcem_module.S_estep(N, Yn, Theta, P, mode
53					#if i <= 3:
54					# print Sns
55					
56					# M-step
57	100	254	2.5	0.0	Theta_old = Theta
58	100	682949	6829.5	0.9	Theta = mcem_module.Theta_mstep(Yn, Sns, model=
59	100	141408	1414.1	0.2	P = mcem_module.P_mstep(Sns)
60					
61					# Store result across iterations
62	100	312	3.1	0.0	Thetas[:, i] = Theta
63	100	306	3.1	0.0	Ps[:, :, i] = P
64					
65					
66					#Stop condition based on convergence
67	100	119	1.2	0.0	if (tol is not None) and (np.allclose(Theta, T
68					break
69					
70	100	120	1.2	0.0	i += 1
71					
72	1	1	1.0	0.0	return Thetas, Ps

- Optimized

Timer unit: 1e-06 s

Total time: 0.029493 s

File: changepoint/mcem\_opt.py

Function: theta\_mstep at line 15

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					

15					def theta_mstep(k, Yn, Sns, model):
16					'''
17					Calculate theta_k, using N samples of Sn. Vectorize
18					'''
19	200	17869	89.3	60.6	Uks = np.sum(Yn[:, np.newaxis] * (Sns == k), axis=0)
20	200	9773	48.9	33.1	Nks = np.sum(Sns == k, axis=0)
21					
22	200	159	0.8	0.5	if model == "binary":
23					return 1.0 * Uks.sum() / Nks.sum()
24	200	97	0.5	0.3	if model == "poisson":
25	200	1595	8.0	5.4	return 1.0 * Uks.sum() / Nks.sum()

Total time: 67.3543 s

File: changepoint/mcem\_sampler.py

Function: mcem\_sampler at line 7

Line #	Hits	Time	Per Hit	% Time	Line Contents
7					def mcem_sampler(Yn, model, m, mcem_module, tol=None):
8					'''
9					Calculate the MLE of Theta and P, using Monte Carlo
10					
11					We use EM because the latent state vector is not o
12					We use Monte Carlo method because the Q function i
13					
14					Args
15					Yn: array, time series data
16					model: string, e.g. "binary", "poisson"
17					m: int, number of change points
18					mcem: module, e.g. mcem, mcem_opt
19					Use the regular of optimized version
20					'''
21					# Initialize
22	1	4	4.0	0.0	n = len(Yn) ; m = m
23	1	151	151.0	0.0	P = init.P(m + 1)
24					
25	1	5	5.0	0.0	P[0, 0] = 0.8 # Paper
26	1	2	2.0	0.0	P[0, 1] = 1 - 0.8
27					
28	1	2	2.0	0.0	if (m == 2) and (model == "poisson"):
29					# Without giving a high chance of staying at s
30					P[1, 1] = 0.9
31					P[1, 2] = 1 - 0.9
32					
33					# Initialize Theta
34	1	2	2.0	0.0	if model == "binary":
35					Theta = np.repeat(0.5, m + 1)
36	1	2	2.0	0.0	elif model == "poisson":
37					#Theta = np.repeat(2, m + 1)
38	1	34	34.0	0.0	Theta = 2 + np.random.rand(m + 1)
39					
40					# N is the number of Sn sample
41					# According to paper, start N = 1 and increases ov
42	1	52	52.0	0.0	Ns = np.linspace(1, 300, 10).astype(np.int64)

```

43      1      7      7.0      0.0      Thetas = np.empty((m + 1, 100))
44      1      3      3.0      0.0      Ps = np.empty((m + 1, m + 1, 100))
45
46                                          # Start 100 MCEM steps
47      1      1      1.0      0.0      i = 0
48     101     125      1.2      0.0      while i < 100:
49     100     184      1.8      0.0          N = Ns[i / 10]
50
51                                          # E-step
52     100    67308444  673084.4      99.9      Sns = mcm_module.S_estep(N, Yn, Theta, P, model=
53                                          #if i <= 3:
54                                          #    print Sns
55
56                                          # M-step
57     100      270      2.7      0.0      Theta_old = Theta
58     100    34774     347.7      0.1      Theta = mcm_module.Theta_mstep(Yn, Sns, model=
59     100     9408     94.1      0.0      P = mcm_module.P_mstep(Sns)
60
61                                          # Store result across iterations
62     100      309      3.1      0.0      Thetas[:, i] = Theta
63     100      294      2.9      0.0      Ps[:, :, i] = P
64
65
66                                          #Stop condition based on convergence
67     100      120      1.2      0.0      if (tol is not None) and (np.allclose(Theta, T
68                                          break
69
70     100      124      1.2      0.0      i += 1
71
72      1      1      1.0      0.0      return Thetas, Ps

```

## 8 Comparison with existing packages

We compare our algorithm with:

- The R implementation in package **MCMCpack**, which produces identical result to ours (See Figures below). (In fact, the marginal likelihood calculation is more similar to ours than to the original paper).
- Another Bayesian change point model, by Barry and Hartigan (1993), which gives similar conclusion. This model detect the number of regimes in the data, their means and variances. The result also suggests one break point, with two regimes having means =  $[3, 1]$ , similar to our Poisson model (See Figures below).

These two packages in R / C++ implementation run a lot faster than our implementation however (2.83s and 1.86s, compared to  $> 1$  min in our implementation).

1 loops, best of 3: 2.82 s per loop

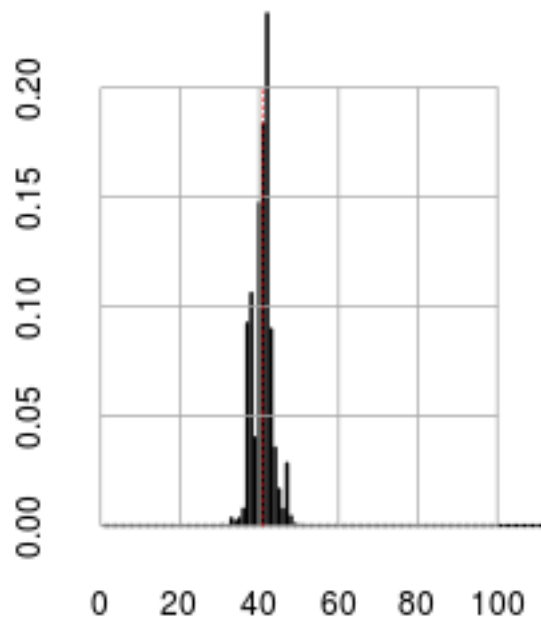
1 loops, best of 3: 1.86 s per loop

### 8.0.1 MCMCpack implementation of Chib (1998)


```
[1] "log-likelihood -172.194701317687"
```

```
[1] "log-marginal-likelihood -178.354928104687"
```

**Posterior Density of Regime Change Probability**



**Posterior Regime Probability**



### 8.0.2 bcp Implementation of Barry and Hartigan (1993)

